

Лабораторная работа №2 по
дисциплине
«Методы машинного обучения» на
тему
«Изучение библиотек обработки данных»

Выполнил:
студент группы ИУ5-21М

Наинг Ко Ко Линн

Москва — 2020 г.

1. Цель лабораторной работы

Изучить библиотеки обработки данных Pandas и PandaSQL [1].

2. Задание

Задание состоит из двух частей [1].

2.1. Часть 1

Требуется выполнить первое демонстрационное задание под названием «Exploratory data analysis with Pandas» со страницы курса mlcourse.ai.

2.2. Часть 2

Требуется выполнить следующие запросы с использованием двух различных библиотек — Pandas и PandaSQL:

- один произвольный запрос на соединение двух наборов данных,
- один произвольный запрос на группировку набора данных с использованием функций агрегирования.

Также требуется сравнить время выполнения каждого запроса в Pandas и PandaSQL.

3. Ход выполнения работы

3.1. Часть 1

Ниже приведён демонстрационный Jupyter-ноутбук «Exploratory data analysis with Pandas» курса mlcourse.ai (файл `assignment01_pandas_uci_adult.ipynb`). Все пояснения приведены на исходном языке ноутбука — на английском.



mlcourse.ai – Open Machine Learning Course

Author: Yury Kashnitskiy. Translated and edited by Sergey Isaev, Artem Trunov, Anastasia Manokhina, and Yuanyuan Pao. This material is subject to the terms and conditions of the Creative Commons CC BY-NC-SA 4.0 license. Free use is permitted for any noncommercial purpose.

Assignment #1 (demo) Exploratory data analysis with Pandas

In this task you should use Pandas to answer a few questions about the Adult dataset.

Unique values of all features (for more information, please see the links above):

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US (Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.
- salary: >50K, <=50K.

Importing all required packages:

```
In [1]: import pandas as pd
```

Setting maximum display width for text report [2]:

```
In [2]: pd.set_option("display.width", 70)
```

Loading data:

```
In [3]: data = pd.read_csv('adult11.csv')
```

```
data.head()
```

```
1 data = pd.read_csv('adult11.csv')
2 data.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week	native-country	salary
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0.0	0.0	40.0	United-States	<=50K
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0.0	0.0	50.0	United-States	<=50K
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0.0	0.0	40.0	United-States	>50K
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688.0	0.0	40.0	United-States	>50K
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0.0	0.0	30.0	United-States	<=50K

1. How many men and women (sex feature) are represented in this dataset?

```
In [4]: data["gender"].value_counts()
```

```
1 data["gender"].value_counts()
```

```
Male      19301
Female     9553
Name: gender, dtype: int64
```

2. What is the average age (age feature) of women?

```
In [5]: data[data["gender"] == "Female"]["age"].mean()
```

```
1 data[data["gender"] == "Female"]["age"].mean()
```

```
36.98230922223385
```

3. What is the percentage of German citizens (native-country feature)?

```
In [6]: print("{0:%}".format(data[data["native-country"] == "Germany"]
.shape[0] / data.shape[0]))
```

```
1 print("{0:%}".format(data[data["native-country"] == "Germany"]
2 .shape[0] / data.shape[0]))
```

```
0.440132%
```

4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature) and those who earn less than 50K per year?

```
In [7]: ages1 = data[data["salary"] == "<=50K"]["age"]

ages2 = data[data["salary"] == ">50K"]["age"]
print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))
print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))
```

```
1 ages1 = data[data["salary"] == "<=50K"]["age"]
2 ages2 = data[data["salary"] == ">50K"]["age"]
3 print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))
4 print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))
```

```
<=50K: = 36.90999590890495 ± 14.166242134563927 years
>50K: = 44.1671772428884 ± 10.52773245753095 years
```

Is it true that people who earn more than 50K have at least high schooleducation? (education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)

```
In [8] high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",
                             "Assoc-voc", "Masters", "Doctorate"])

def high_educated(e):
    return e in high_educations
data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

```
1 high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",
2 | | | | | | | | | | "Assoc-voc", "Masters", "Doctorate"])
3 def high_educated(e):
4 |     return e in high_educations
5 data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

```
False
```

Display age statistics for each race (race feature) and each gender (sex feature). Use groupby() and describe(). Find the maximum age of men of Amer-Indian-Eskimo race.

```
In [9]: data.groupby(["race", "sex"])["age"].describe()
```

```
1 data.groupby(["race", "gender"])["age"].describe()
```



		count	mean	std	min	25%	50%	75%	max
	race	gender							
	Amer-Indian-Eskimo	Female	115.0	34.930435	12.322937	17.0	25.00	33.0	63.0
		Male	160.0	37.125000	11.781955	17.0	29.00	35.0	82.0
	Asian-Pac-Islander	Female	300.0	36.350000	13.155139	17.0	26.00	35.0	81.0
		Male	564.0	38.843972	12.891560	18.0	29.00	37.0	90.0
	Black	Female	1338.0	38.354260	12.967195	17.0	28.00	37.0	81.0
		Male	1422.0	38.059775	13.239203	17.0	28.00	36.0	90.0
	Other	Female	80.0	30.737500	10.861317	17.0	22.75	28.0	64.0
		Male	153.0	34.954248	11.800613	17.0	26.00	32.0	67.0
	White	Female	7720.0	36.864378	14.488046	17.0	25.00	35.0	90.0
		Male	17002.0	39.650629	13.475728	17.0	29.00	38.0	90.0

```
In [10]: data[(data["race"] == "Amer-Indian-Eskimo")
& (data["gender"] == "Male")]["age"].max()
```

```
1 data[(data["race"] == "Amer-Indian-Eskimo")
2 & (data["gender"] == "Male")]["age"].max()
```

```
82
```

8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)? Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Marriedspouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
In [11]: def is_married(m):
        return m.startswith("Married")
data["married"] = data["marital-status"].map(is_married)
(data[(data["gender"] == "Male") & (data["salary"] == ">50K")][
    "married"].value_counts())
```

```
1 def is_married(m):
2     return m.startswith("Married")
3 data["married"] = data["marital-status"].map(is_married)
4 (data[(data["gender"] == "Male") & (data["salary"] == ">50K")][
5     "married"].value_counts())
```

```
True      5237
False      567
Name: married, dtype: int64
```

9. What is the maximum number of hours a person works per week (hours-per-week feature)? How many people work such a number of hours, and what is the percentage of those who earn a lot (>50K) among them?

```
In [12]: m = data["hours-per-week"].max()

print("Maximum is {} hours/week.".format(m))
people = data[data["hours-per-week"] == m]
c = people.shape[0]
print("{} people work this time at week.".format(c))
s = people[people["salary"] == ">50K"].shape[0]
print("{} get >50K salary.".format(s / c))
```

```

1 m = data["hours-per-week"].max()
2 print("Maximum is {} hours/week.".format(m))
3 people = data[data["hours-per-week"] == m]
4 c = people.shape[0]
5 print("{} people work this time at week.".format(c))
6 s = people[people["salary"] == ">50K"].shape[0]
7 print("{} get >50K salary.".format(s / c))

```

```

Maximum is 99.0 hours/week.
75 people work this time at week.
32.000000% get >50K salary.

```

10. Count the average time of work (hours-per-week) for those who earn a little and a lot (salary) for each country (native-country). What will these be for Japan?

```

In [13]: p = pd.crosstab(data["native-country"], data["salary"],
values=data['hours-per-week'], aggfunc="mean")p
p

```

```

1 p = pd.crosstab(data["native-country"], data["salary"],
2 values=data['hours-per-week'], aggfunc="mean")p
3 p

```

	salary	<=50K	>50K
native-country			
?		38.637883	45.384058
Cambodia		41.307692	47.000000
Canada		36.411765	46.456522
China		35.900000	44.380952
Columbia		40.513514	56.250000
Cuba		40.825397	40.937500
Dominican-Republic		40.962963	42.800000
Ecuador		36.875000	47.333333
El-Salvador		35.231707	43.600000
England		38.148936	46.384615
France		41.100000	39.166667
Germany		38.223404	45.969697
Greece		44.333333	56.428571
Guatemala		38.062500	40.000000

Haiti	34.386364	39.166667
Honduras	32.090909	50.000000
Hong	38.181818	40.000000
Hungary	36.000000	42.666667
India	39.207547	44.470588
Iran	39.458333	48.000000
Ireland	40.933333	44.111111
Italy	38.333333	44.500000
Jamaica	39.780000	43.909091
Japan	38.972973	43.666667
Laos	37.777778	NaN

Mexico	39.870229	46.214286
Nicaragua	36.760000	40.000000
Outlying-US(Guam-USVI-etc)	41.666667	40.000000
Peru	37.318182	40.000000
Philippines	39.096000	44.585366
Poland	38.795455	38.300000
Portugal	40.000000	47.100000
Puerto-Rico	39.101852	39.866667
Scotland	40.000000	NaN
South	42.721311	51.666667
Taiwan	39.600000	44.000000
Thailand	46.500000	53.750000
Trinidad&Tobago	39.266667	NaN
United-States	38.883042	45.430669
Vietnam	39.523810	40.000000
Yugoslavia	36.500000	32.500000

```
In [14]: p.loc["Japan"]
```

```
1 p.loc["Japan"]
```

```
salary  
<=50K    38.972973  
>50K     43.666667  
Name: Japan, dtype: float64
```

3.2. Часть 2

Импортируем pandasql:

```
In [15]: !pip install pandasql
```

```
from pandasql import sqldf
```

```
import pandas as pd
```

```
In [16]: from pandasql import sqldf
```

```
pysqldf = lambda q: sqldf(q, globals())
```

```
1 from pandasql import sqldf  
2 pysqldf = lambda q: sqldf(q, globals())
```

```
[ ] 1 wind = (pd.read_csv('wind speed.csv', header=None,  
2 names=["row", "UNIX", "date",  
3 "time", "speed", "text"]))  
4 .drop("text", axis=1))  
5 temp = (pd.read_csv('temperature.csv', header=None,  
6 names=["row", "UNIX", "date",  
7 "time", "temperature", "text"]))  
8 .drop("text", axis=1))
```

```
In [17]: wind.head()
```

```
1 wind.head()
```

	row	UNIX	date	time	speed
0	1	1475315718	2016-09-30	23:55:18	7.87
1	2	1475315423	2016-09-30	23:50:23	7.87
2	3	1475315124	2016-09-30	23:45:24	9.00
3	4	1475314821	2016-09-30	23:40:21	13.50
4	5	1475314522	2016-09-30	23:35:22	15.75

```
In [18]: wind.dtypes
```

```
[ ] 1 wind.dtypes
```

```
↳ row      int64
   UNIX     int64
   date     object
   time     object
   speed    float64
   dtype: object
```

```
In [19]: temp.head()
```

```
1 temp.head()
```

```
↳
```

	row	UNIX	date	time	temperature
0	1	1475315718	2016-09-30	23:55:18	48
1	2	1475315423	2016-09-30	23:50:23	48
2	3	1475315124	2016-09-30	23:45:24	48
3	4	1475314821	2016-09-30	23:40:21	48
4	5	1475314522	2016-09-30	23:35:22	48

```
In [20]: temp.dtypes
```

```
[ ] 1 temp.dtypes
```

```
↳ row      int64
   UNIX     int64
   date     object
   time     object
   temperature int64
   dtype: object
```

```
In [21]: wind.merge(temp[["UNIX", "temperature"]], on="UNIX").head()
```

```
1 wind.merge(temp[["UNIX", "temperature"]], on="UNIX").head()
```

```
↳
```

	row	UNIX	date	time	speed	temperature
0	1	1475315718	2016-09-30	23:55:18	7.87	48
1	2	1475315423	2016-09-30	23:50:23	7.87	48
2	3	1475315124	2016-09-30	23:45:24	9.00	48
3	4	1475314821	2016-09-30	23:40:21	13.50	48
4	5	1475314522	2016-09-30	23:35:22	15.75	48

```
In [22]: %%timeit
```

```
wind.merge(temp[["UNIX", "temperature"]], on="UNIX")
```

```
1 %%timeit
2 wind.merge(temp[["UNIX", "temperature"]], on="UNIX")
```

```
100 loops, best of 3: 10.9 ms per loop
```

```
In [23]: pysqldf("""SELECT w.row, w.UNIX, w.date, w.time,
w.speed, t.temperature
FROM wind AS w JOIN temp AS t
ON w.UNIX = t.UNIX
""").head()
```

```
1 pysqldf("""SELECT w.row, w.UNIX, w.date, w.time,
2 w.speed, t.temperature
3 FROM wind AS w JOIN temp AS t
4 ON w.UNIX = t.UNIX
5 """).head()
```

	row	UNIX	date	time	speed	temperature
0	1	1475315718	2016-09-30	23:55:18	7.87	48
1	2	1475315423	2016-09-30	23:50:23	7.87	48
2	3	1475315124	2016-09-30	23:45:24	9.00	48
3	4	1475314821	2016-09-30	23:40:21	13.50	48
4	5	1475314522	2016-09-30	23:35:22	15.75	48

```
In [24]: %%timeit pysqldf("""SELECT w.row,
w.UNIX, w.date, w.time,
w.speed,
t.temperature FROM wind
AS w JOIN temp AS t
ON w.UNIX = t.UNIX
""")
```

```
1 %%timeit
2 pysqldf("""SELECT w.row, w.UNIX, w.date, w.time,
3 w.speed, t.temperature
4 FROM wind AS w JOIN temp AS t
5 ON w.UNIX = t.UNIX
6 """)
```

```
1 loop, best of 3: 505 ms per loop
```

Видно, что pandasql в 50 раз медленнее, чем pandas.

Сгруппируем набор данных с использованием функций агрегирования различными способами:

```
In [25]: wind.groupby("date")["speed"].mean().head()
```

```
1 wind.groupby("date")["speed"].mean().head()
```

date	
2016-09-01	6.396560
2016-09-02	5.804086
2016-09-03	4.960248
2016-09-04	5.184571
2016-09-05	5.830676

Name: speed, dtype: float64

```
In [26]: %%timeit
wind.groupby("date")["speed"].mean()
```

```
[ ] 1 %%timeit
2 wind.groupby("date")["speed"].mean()
```

100 loops, best of 3: 2.72 ms per loop

```
In [27]: pysqldf("""SELECT date, AVG(speed)
FROM wind
GROUP BY date
""").head()
```

```
[ ] 1 pysqldf("""SELECT date, AVG(speed)
2 FROM wind
3 GROUP BY date
4 """).head()
```

	date	AVG(speed)
0	2016-09-01	6.396560
1	2016-09-02	5.804086
2	2016-09-03	4.960248
3	2016-09-04	5.184571
4	2016-09-05	5.830676

```
In [28]: %%timeit
        pysqldf("""SELECT date,
                AVG(speed)
                FROM wind
                GROUP BY date
                """)
```

```
[ ] 1  %%timeit
     2  pysqldf("""SELECT date, AVG(speed)
     3  FROM wind
     4  GROUP BY date
     5  """)
```

↳ 1 loop, best of 3: 202 ms per loop

Здесь разница уже более чем в 100 раз. Таким образом для таких простых запросов проще использовать Pandas.

Список литературы

- [1] Гапанюк Ю. Е. Лабораторная работа «Изучение библиотек обработки данных» [Электронный ресурс] // GitHub. — 2019. — Режим доступа: https://github.com/ugapanyuk/ml_course/wiki/LAB_PANDAS (дата обращения: 20.02.2019).
- [2] pandas 0.24.1 documentation [Electronic resource] // PyData. — 2019. — Access mode: <http://pandas.pydata.org/pandas-docs/stable/> (online; accessed: 20.02.2019).
- [3] You are my Sunshine [Electronic resource] // Space Apps Challenge. — 2017. — Access mode: <https://2017.spaceappschallenge.org/challenges/earth-and-us/you-are-my-sunshine/details> (online; accessed: 22.02.2019).
- [4] yhat/pandasql: sqldf for pandas [Electronic resource] // GitHub. — 2017. — Access mode: <https://github.com/yhat/pandasql> (online; accessed: 22.02.2019).
- [5] Team The IPython Development. IPython 7.3.0 Documentation [Electronic resource] // Read the Docs. — 2019. — Access mode: <https://ipython.readthedocs.io/en/stable/> (online; accessed: 20.02.2019).